

GAMs

Identifiability
Effective Degrees
of Freedom
Smoothing
Parameter
Selection

Software

Model
Fitting/Plotting
Model Plotting

In-Class Exercises

BIOS 7720: Applied Functional Data Analysis

Lecture 3: Generalized Additive Models (GAMs) Part 1

Andrew Leroux

March 9, 2021

GAMs

Identifiability
Effective Degrees
of Freedom
Smoothing
Parameter
Selection

Software

Model
Fitting/Plotting
Model Plotting

In-Class Exercises

- Homework 1 due date now 3/19
- Other homework 1 things
- Lecture recordings
- Best form of communication?
- Office hours start this week
 - One off this week: 5/12 at 10:00AM
 - Hereafter: Fridays at 9:00AM

Roadmap

GAMs

Identifiability
Effective Degrees
of Freedom
Smoothing
Parameter
Selection

Software

Model
Fitting/Plotting
Model Plotting

In-Class Exercises

- 1 Introduce GAMs and associated concepts
- 2 In-class exercises

Introduction

GAMs

Identifiability
Effective Degrees
of Freedom
Smoothing
Parameter
Selection

Software

Model
Fitting/Plotting
Model Plotting

In-Class Exercises

- General form

$$y_i \sim \text{EF}(\mu_i, \sigma)$$
$$g(E[y_i | \mathbf{x}_i]) = f(\mathbf{x}_i)$$

- An example with $g(x) = x$ and $\mathbf{x}_i = [x_{i1}, x_{i2}, x_{i3}, x_{i4}]^t$

$$E[y_i | \mathbf{x}_i] = f(x_{i1}, x_{i2}, x_{i3}, x_{i4})$$
$$= f_1(x_{i1}) + f_2(x_{i2}) + f_3(x_{i3}, x_{i4})$$

Identifiability

GAMs

Identifiability

Effective Degrees
of Freedom

Smoothing
Parameter
Selection

Software

Model
Fitting/Plotting
Model Plotting

In-Class Exercises

- Up until now we've implicitly assumed the $c1_{N \times 1} \in \text{span}(\Phi)$ where $c \in \mathbb{R}^1$
- Suppose we want to fit the model

$$\begin{aligned}y_i &= \alpha_0 + f(x_i) \\ &= \alpha_0 + \Phi\xi\end{aligned}$$

- This model is not identifiable since

$$y_i = (\alpha_0 + b) + (\Phi\xi - b)$$

- Need to impose some constraint

Identifiability in *mgcv*

GAMs

Identifiability

Effective Degrees
of Freedom

Smoothing
Parameter
Selection

Software

Model
Fitting/Plotting
Model Plotting

In-Class Exercises

- Consider the constraint $\sum_{i=1}^N f(x_i) = 0$

$$\sum_{i=1}^N f(x_i) = 0$$

$$\mathbf{1}^t \Phi \xi = 0$$

- For this to be true for all ξ , we require $\mathbf{1}^t \Phi = 0$
- Easy solution, center Φ columnwise

Identifiability in *mgcv*

GAMs

Identifiability

Effective Degrees
of Freedom

Smoothing
Parameter
Selection

Software

Model
Fitting/Plotting
Model Plotting

In-Class
Exercises

- Column centered basis matrix: $\tilde{\Phi} = (I_{N \times N} - N^{-1}\mathbf{1}\mathbf{1}^t)\Phi$
- Problem: $\tilde{\Phi}$ is not full rank (homework problem)
- Solution:
 - Delete a column from $\tilde{\Phi}$
 - Delete the corresponding row/column from S

Identifiability in *mgcv*

GAMs

Identifiability

Effective Degrees
of Freedom

Smoothing
Parameter
Selection

Software

Model
Fitting/Plotting
Model Plotting

In-Class Exercises

$$y_i = \log(1 + x_i^2) + \epsilon_i$$

$$x_i \sim N(0, \sigma_x^2)$$

$$\epsilon_i \sim N(0, \sigma_\epsilon^2)$$

```
library("mgcv")
set.seed(-909)
N <- 1000
f <- function(x) log(1+x^2)
x <- rnorm(N, sd=3)
y <- f(x) + rnorm(N, mean=0, sd=0.5)
```


Identifiability in *mgcv*

What if we don't impose the identifiability constraint?

```
## number of basis functions to use in fitting
K <- 10

## sequence of x values to predict on (for plotting)
xind <- seq(min(x), max(x), len=1000)

## set up the unconstrained smooth, apply identifiability constraints
sm <- smoothCon(s(x, bs="cr", k=K), data=data.frame(x=x))[[1]]
Phi<- sm$X

## fit using lm()
fit_1 <- lm(y~Phi)

## get the design matrix for our initial fit
lp_1 <- PredictMat(sm, data=data.frame(x=xind))

## get the estimated function at our x-values for predicting
fhat_1 <- lp_1 %*% coef(fit_1)

## Error in lp_1 %*% coef(fit_1): non-conformable arguments
```

Identifiability in *mgcv*

GAMs

Identifiability

Effective Degrees
of Freedom

Smoothing
Parameter
Selection

Software

Model
Fitting/Plotting
Model Plotting

In-Class Exercises

```
## number of basis functions to use in fitting
K <- 10
## sequence of x values to predict on (for plotting)
xind <- seq(min(x), max(x), len=1000)
## set up the unconstrained smooth, apply identifiability constraints
sm <- smoothCon(s(x, bs="cr", k=K), data=data.frame(x=x))[[1]]
Phi <- sm$X
Phi_mn <- colMeans(Phi)
Phi_tilde <- sweep(Phi, 2, Phi_mn)
Phi_tilde <- Phi_tilde[, -K]
## fit using lm() -- no need to suppress the intercept
fit_1 <- lm(y~Phi_tilde)
## get the design matrix for our initial fit
lp_1 <- PredictMat(sm, data=data.frame(x=xind))
lp_1 <- cbind(1, sweep(lp_1, 2, Phi_mn)[, -K])
## get the estimated function at our x-values for predicting
fhat_1 <- lp_1 %*% coef(fit_1)
```

Identifiability in *mgcv*

GAMs

Identifiability

Effective Degrees
of Freedom

Smoothing
Parameter
Selection

Software


Model
Fitting/Plotting
Model Plotting

In-Class
Exercises

```
## do the fit with the identifiability constraint absorbed
sm_cons <- smoothCon(s(x, bs="cr", k=K), data=data.frame(x=x),
                     absorb.cons = TRUE)[[1]]
fit_2 <- lm(y~sm_cons$X)
lp_2 <- cbind(1, PredictMat(sm_cons, data=data.frame(x=xind)))
fhat_2 <- lp_2 %*% coef(fit_2)
```

Identifiability in *mgcv*

```
matplot(xind, cbind(fhat_1, fhat_2, f(xind)),type='l',xlab="x",ylab="y",lwd=2)  
legend("top",c("Manually Constrain", "Use mgcv","Truth"), col=1:3,lty=1:3,bty='n',
```



GAMs

Identifiability

Effective Degrees
of Freedom

Smoothing
Parameter
Selection

Software

Model
Fitting/Plotting
Model Plotting

In-Class Exercises

Identifiability in *mgcv*

GAMs

Identifiability

Effective Degrees of Freedom

Smoothing
Parameter
Selection


Software

Model
Fitting/Plotting
Model Plotting


In-Class
Exercises

```
layout(matrix(c(1,2),1,2))  
matplot(xind,lp_1,type='l',main="Design Matrix: Manually Constrain",xlab="x")  
matplot(xind,lp_2,type='l',main="Design Matrix: mgcv Constrain",xlab="x")
```

Design Matrix: Manually Constrain



Design Matrix: mgcv Constrain



Effective Degrees of Freedom

GAMs

Identifiability

Effective Degrees of Freedom

Smoothing Parameter Selection

Software

Model Fitting/Plotting Model Plotting

In-Class Exercises

- Recall from linear regression $\hat{\sigma}^2 = ||\mathbf{y} - \hat{\mathbf{y}}||^2 / (N - p)$
- With penalized splines, we want large p (flexibility)
- But what if the actual $\hat{f}(x)$ doesn't use all p df?
- as $\lambda \rightarrow \infty$, $\hat{f}(x) = \beta x$, so 1 df "used"

Effective Degrees of Freedom

GAMs

Identifiability

Effective Degrees of Freedom

Smoothing Parameter Selection

Software

Model Fitting/Plotting Model Plotting

In-Class Exercises

- Recall from linear regression $\hat{\sigma}^2 = ||\mathbf{y} - \hat{\mathbf{y}}||^2 / (N - p)$
- With penalized splines, we want large p (flexibility)
- But what if the actual $\hat{f}(x)$ doesn't use all p df?
- as $\lambda \rightarrow \infty$, $\hat{f}(x) = \beta x$, so 1 df "used"

Effective Degrees of Freedom

GAMs

Identifiability

Effective Degrees of Freedom

Smoothing Parameter Selection

Software

Model Fitting/Plotting Model Plotting

In-Class Exercises

- Use \mathbf{X} as general notation for our design matrix (e.g. Φ)
- Recall from linear regression $\hat{\mathbf{y}} = \mathbf{H}\mathbf{y}$
- $\mathbf{H} = \mathbf{X}(\mathbf{X}^t\mathbf{X})^{-1}\mathbf{X}^t$ is the matrix that projects \mathbf{y} onto the column space of \mathbf{X}
- $\text{trace}(\mathbf{H}) = P$ (assuming $\mathbf{X}_{N \times P}$ is full column rank)
- Define for a fixed λ : $\text{EDF} = \text{trace}(\mathbf{X}(\mathbf{X}^t\mathbf{X} + \lambda\mathbf{S})^{-1}\mathbf{X}^t)$

Smoothing Parameter Selection

GAMs

Identifiability

Effective Degrees
of Freedom

**Smoothing
Parameter
Selection**

Software

Model
Fitting/Plotting

Model Plotting

In-Class
Exercises

- Ordinary cross-validation
- Generalized cross-validation
- (Marginal) likelihood approaches

Smoothing Parameter Selection: OCV vs GCV

GAMs

Identifiability
Effective Degrees
of Freedom

Smoothing
Parameter
Selection

Software

Model
Fitting/Plotting
Model Plotting

In-Class Exercises

- Ordinary cross-validation is
 - Not defined for general exponential family outcomes
 - Not rotationally invariant
- Solution: generalized cross-validation
- For Gaussian data with identity link

$$\mathcal{V}_0 = N^{-1} \sum_{i=1}^N (y_i - \hat{y}_i)^2 / (1 - H_{ii})^2$$

$$\mathcal{V}_g = N \sum_{i=1}^N (y_i - \hat{y}_i)^2 / (N - \text{trace}(H))^2$$

Smoothing Parameter Selection: OCV vs GCV

GAMs

Identifiability

Effective Degrees of Freedom

Smoothing Parameter Selection

Software

Model Fitting/Plotting

Model Plotting

In-Class Exercises

```
## simulate data  $y_i = 3x_i + \epsilon_i$ 
set.seed(-304); N <- 100; x <- rnorm(N); y <- 3*x + rnorm(N)
## rotate the response and design matrix using an orthogonal matrix
Q <- qr.Q(qr(cbind(1,x)),complete=TRUE)
Qy <- Q %*% y
Qx <- Q %*% cbind(1,x)
## fit the two models
fit_r <- lm(Qy ~ Qx - 1)
fit <- lm(y ~ x)
cbind(coef(fit_r), coef(fit))

##           [,1]      [,2]
## Qx  0.1678198 0.1678198
## Qxx 2.9509341 2.9509341

## GCV and OCV for the two fits
OCV_r <- mean(residuals(fit_r)^2/(1-influence(fit_r)$hat))
OCV <- mean(residuals(fit)^2/(1-influence(fit)$hat))
GCV_r <- N*sum(residuals(fit_r)^2)/(N-sum(influence(fit_r)$hat))^2
GCV <- N*sum(residuals(fit)^2)/(N-sum(influence(fit)$hat))^2
data.frame("Model" = c("Original","Rotated"), GCV=c(GCV,GCV_r), OCV=c(OCV,OCV_r))

##      Model      GCV      OCV
## 1 Original 0.9206058 0.9023693
## 2 Rotated 0.9206058 0.8951349
```

Smoothing Parameter Selection: GCV

GAMs

Identifiability
Effective Degrees
of Freedom

Smoothing
Parameter
Selection

Software

Model
Fitting/Plotting
Model Plotting

In-Class Exercises

- Rotation by an orthogonal matrix Q
 - Yields identical inference on regression coefficients
 - Different OCV criteria
 - Identical GCV criteria
- GCV exists for generalized outcomes using Deviance
- In practice GCV tends to undersmooth at realistic sample sizes
- Can introduce a tuning parameter $\gamma > 1$ to increase smoothness

$$\mathcal{V}_g = N \sum_{i=1}^N (y_i - \hat{y}_i)^2 / (N - \gamma \text{trace}(H))^2$$

- Not clear how to choose γ in a principled way

Smoothing Parameter Selection: likelihood approaches

GAMs

Identifiability

Effective Degrees of Freedom

Smoothing Parameter Selection

Software

Model Fitting/Plotting

Model Plotting

In-Class Exercises

- As an alternative to GCV we can use likelihood based methods
- λ can be estimated as a variance parameter
 - Allows us to frame penalized splines as mixed models
 - More details next class
- Both ML and REML approaches for exponential family responses
- REML is considered the “default” by many for applied work

Simulate Gaussian and Binary Data

GAMs

Identifiability
Effective Degrees
of Freedom
Smoothing
Parameter
Selection

Software

Model
Fitting/Plotting
Model Plotting

In-Class
Exercises

- Simulate data according to the model

$$\begin{aligned}g(E[y_i | \mathbf{x}_i]) &= \eta_i \\&= 1 + f_1(x_{i1}) + f_2(x_{i2}) + f_3(x_{i3}) + f_4(x_{i4}) \\&= 1 + \sin(\pi x_{i1}) + 2x_{i2} + (x_{i3} + x_{i3}^2) + \cos(\pi x_{i4})x_{i4}\end{aligned}$$

- Where y_i is either Gaussian or Binary (logit link)
- Goals
 - Estimate these models using GCV and REML
 - Quickly plot the results
 - Get point/variance estimates directly

Simulate Gaussian and Binary Data

GAMs

- Identifiability
- Effective Degrees of Freedom
- Smoothing Parameter Selection

Software

- Model Fitting/Plotting
- Model Plotting

In-Class Exercises

```
library("mgcv"); set.seed(80487); N <- 200
P <- 4
X <- matrix(rnorm(N*P), N, P)
## set up the association structures
f1 <- function(x) sin(pi*x)
f2 <- function(x) 2*x
f3 <- function(x) 0.25*x^3
f4 <- function(x) cos(pi*x)*x
## get the linear predictor
eta <- 1 + f1(X[,1]) + f2(X[,2]) + f3(X[,3]) + f4(X[,4])
## simulate gaussian outcomes
y_g <- eta + rnorm(N, sd=1)
## simulate binary outcomes
pr_y <- 1/(1+exp(-eta))
y_b <- vapply(pr_y, function(x) sample(c(0,1), size=1, prob=c(1-x,x)),
             numeric(1))
```

Fit the Models using `mgcv::gam()`

GAMs

Identifiability
Effective Degrees
of Freedom
Smoothing
Parameter
Selection

Software

Model
Fitting/Plotting
Model Plotting

In-Class
Exercises

```
## combine data into a dataframe
df_fit <- data.frame(y_g=y_g, y_b=y_b, X)
## set up basis type for all smooth terms
bs <- "cr"
## number of basis functions for all smooth terms
K <- 20
## fit the models on the Gaussian data
fit_g_GCV <- gam(y_g ~ s(X1, bs=bs, k=K) + s(X2, bs=bs, k=K) +
                  s(X3, bs=bs, k=K) + s(X4, bs=bs, k=K),
                  family=gaussian(), method="GCV.Cp", data=df_fit)
fit_g_REML <- gam(y_g ~ s(X1, bs=bs, k=K) + s(X2, bs=bs, k=K) +
                  s(X3, bs=bs, k=K) + s(X4, bs=bs, k=K),
                  family=gaussian(), method="REML", data=df_fit)
## fit the models on the binary data
fit_b_GCV <- gam(y_b ~ s(X1, bs=bs, k=K) + s(X2, bs=bs, k=K) +
                  s(X3, bs=bs, k=K) + s(X4, bs=bs, k=K),
                  family=binomial(), method="GCV.Cp", data=df_fit)
fit_b_REML <- gam(y_b ~ s(X1, bs=bs, k=K) + s(X2, bs=bs, k=K) +
                  s(X3, bs=bs, k=K) + s(X4, bs=bs, k=K),
                  family=binomial(), method="REML", data=df_fit)
```


Plotting the results *mgcv::gam()*

GAMs

Identifiability
Effective Degrees
of Freedom
Smoothing
Parameter
Selection

Software

Model
Fitting/Plotting
Model Plotting

In-Class
Exercises

- *mgcv* has default plotting methods
- just call `plot(fit)` on your fitted model. For example:

```
par(mfrow=c(2,2))  
plot(fit_g_GCV, shade=TRUE)
```

- See `?plot.gam` for additional options
- Can gain finer control by looping over coefficients

Plotting the results `mgcv::gam()`: Gaussian, GCV

GAMs

Identifiability
Effective Degrees
of Freedom
Smoothing
Parameter
Selection

Software

Model
Fitting/Plotting
Model Plotting

In-Class
Exercises

```
par(mfrow=c(2,2))
nx_pred <- 1000
xind_pred <-
  lapply(1:P, function(x){
    rn_x <- range(X[,x])
    seq(rn_x[1], rn_x[2], len=nx_pred)
  })
fn_ls <- list(f1,f2,f3,f4)
for(p in 1:P){
  plot(fit_g_GCV, select=p, shade=TRUE)
  lines(xind_pred[[p]], fn_ls[[p]](xind_pred[[p]]), col='red', lwd=2, lty=2)
}
```

Plotting the results: Gaussian, GCV





GAMs

Identifiability
Effective Degrees
of Freedom
Smoothing
Parameter
Selection

Software

Model
Fitting/Plotting
Model Plotting

In-Class
Exercises



Plotting the results: Gaussian, REML





GAMs

Identifiability
Effective Degrees
of Freedom
Smoothing
Parameter
Selection

Software

Model
Fitting/Plotting
Model Plotting

In-Class
Exercises



Plotting the results: Binary, GCV





GAMs

Identifiability
Effective Degrees
of Freedom
Smoothing
Parameter
Selection

Software

Model
Fitting/Plotting
Model Plotting

In-Class
Exercises



Plotting the results: Binary, REML





GAMs

Identifiability
Effective Degrees
of Freedom
Smoothing
Parameter
Selection

Software

Model
Fitting/Plotting
Model Plotting

In-Class
Exercises



Extracting These Quantities

GAMs

Identifiability
Effective Degrees
of Freedom
Smoothing
Parameter
Selection

Software

Model
Fitting/Plotting
Model Plotting

In-Class
Exercises

- That's all helpful for a quick loop, but how can we get estimates manually?
- Use *predict.gam()*
- You can obtain predictions on
 - Response scale (type="response")
 - Linear predictor (type="link")
 - Component of the linear predictor (type="terms")
 - "design matrix" (type="lpmatrix")
- Set the argument `se.fit=TRUE` to get component-wise standard errors

Extracting These Quantities

GAMs

Identifiability
Effective Degrees
of Freedom
Smoothing
Parameter
Selection

Software

Model
Fitting/Plotting
Model Plotting

In-Class
Exercises

```
## set up a new data frame with all "X" predictors at a new range of values
xind_pred <- seq(-3,3,len=1000)
df_pred <- data.frame(X1=xind_pred, X2=xind_pred, X3=xind_pred, X4=xind_pred)
head(df_pred)
```

##	X1	X2	X3	X4
## 1	-3.000000	-3.000000	-3.000000	-3.000000
## 2	-2.993994	-2.993994	-2.993994	-2.993994
## 3	-2.987988	-2.987988	-2.987988	-2.987988
## 4	-2.981982	-2.981982	-2.981982	-2.981982
## 5	-2.975976	-2.975976	-2.975976	-2.975976
## 6	-2.969970	-2.969970	-2.969970	-2.969970

```
## get the predicted values at these values for each "type"
yhat_g_REML <- predict(fit_g_REML, newdata=df_pred, type="response", se.fit=TRUE)
etahat_g_REML <- predict(fit_g_REML, newdata=df_pred, type="link", se.fit=TRUE)
smhat_g_REML <- predict(fit_g_REML, newdata=df_pred, type="terms", se.fit=TRUE)
Phi_g_REML <- predict(fit_g_REML, newdata=df_pred, type="lpmatrix", se.fit=TRUE)
```


Extracting These Quantities: $\hat{y}_i, \hat{\eta}_i$

GAMs

Identifiability
Effective Degrees
of Freedom
Smoothing
Parameter
Selection

Software

Model
Fitting/Plotting

Model Plotting

In-Class
Exercises

```
str(yhat_g_REML)

## List of 2
## $ fit : num [1:1000(1d)] -12.3 -12.2 -12.2 -12.1 -12.1 ...
## .. attr(*, "dimnames")=List of 1
## .. ..$ : chr [1:1000] "1" "2" "3" "4" ...
## $ se.fit: num [1:1000(1d)] 1.71 1.7 1.69 1.68 1.67 ...
## .. attr(*, "dimnames")=List of 1
## .. ..$ : chr [1:1000] "1" "2" "3" "4" ...

str(etahat_g_REML)

## List of 2
## $ fit : num [1:1000(1d)] -12.3 -12.2 -12.2 -12.1 -12.1 ...
## .. attr(*, "dimnames")=List of 1
## .. ..$ : chr [1:1000] "1" "2" "3" "4" ...
## $ se.fit: num [1:1000(1d)] 1.71 1.7 1.69 1.68 1.67 ...
## .. attr(*, "dimnames")=List of 1
## .. ..$ : chr [1:1000] "1" "2" "3" "4" ...
```

Are these quantities useful?

Extracting These Quantities: \hat{f}

GAMs

Identifiability
Effective Degrees
of Freedom
Smoothing
Parameter
Selection

Software

Model
Fitting/Plotting
Model Plotting

In-Class
Exercises

```
str(smhat_g_REML)

## List of 2
## $ fit : num [1:1000, 1:4] -0.582 -0.576 -0.571 -0.565 -0.559 ...
## .. attr(*, "dimnames")=List of 2
## .. ..$ : chr [1:1000] "1" "2" "3" "4" ...
## .. ..$ : chr [1:4] "s(X1)" "s(X2)" "s(X3)" "s(X4)"
## $ se.fit: num [1:1000, 1:4] 0.809 0.804 0.798 0.792 0.786 ...
## .. attr(*, "dimnames")=List of 2
## .. ..$ : chr [1:1000] "1" "2" "3" "4" ...
## .. ..$ : chr [1:4] "s(X1)" "s(X2)" "s(X3)" "s(X4)"
## - attr(*, "constant")= Named num 0.936
## .. attr(*, "names")= chr "(Intercept)"

# matplot(smhat_g_REML, type='l', xlab="x", ylab=expression(hat(f)(x)))
```

- Returns a list with two elements
- The estimated smooths \hat{f}_p and corresponding standard errors evaluated at each point in `df_pred`

Extracting These Quantities: Φ

```
str(Phi_g_REML)
```

```
## num [1:1000, 1:77] 1 1 1 1 1 1 1 1 1 1 ...
## - attr(*, "dimnames")=List of 2
## ..$ : chr [1:1000] "1" "2" "3" "4" ...
## ..$ : chr [1:77] "(Intercept)" "s(X1).1" "s(X1).2" "s(X1).3" ...
## - attr(*, "model.offset")= num 0
```

- Let $\mathbf{X}_{\text{new}} = [\mathbf{x}_{1,\text{new}}, \dots, \mathbf{x}_{4,\text{new}}]$ be the matrix of "new" predictor values we want to evaluate at
- Then $\Phi_{\text{new}} = [\Phi_1(\mathbf{x}_{1,\text{new}}), \dots, \Phi_4(\mathbf{x}_{4,\text{new}})]$ is the associated design matrix
- $\hat{\mathbf{y}}_{\text{new}} = \Phi_{\text{new}} \boldsymbol{\theta}$
- Where $\boldsymbol{\theta} = [\beta_0, \hat{\xi}_1, \dots, \hat{\xi}_4]^t$
- Why doesn't *mgcv* return standard errors?

Exercises

GAMs

Identifiability
Effective Degrees
of Freedom
Smoothing
Parameter
Selection

Software

Model
Fitting/Plotting
Model Plotting

In-Class Exercises

- 1 Simulate an additive model where the outcome is Poisson
- 2 Plot the estimated coefficients ($\hat{f}(x)$) with 95% CIs ($\hat{f}(x) \pm 2SE(\hat{f}(x))$)
- 3 Obtain estimates of $\hat{f}(x_{\text{new}})$ for a new set of x values by:
 - a. directly using `predict.gam()` with `type="terms"`
 - b. post multiplying Φ_{new} with θ using the columns/entries associated with each smooth term

Next Class

GAMs

Identifiability
Effective Degrees
of Freedom
Smoothing
Parameter
Selection

Software

Model
Fitting/Plotting
Model Plotting

In-Class Exercises

- (Marginal) likelihood approach to smoothing parameter selection
- Smooths of more than one variable
- Linear functionals
- Interpreting summary output from `mgcv::gam()`